



Use Case Risk Recategorization

Ryef Taimur, Ibrahim Murtaza

1 Purpose

Company and client information is regularly sent to commercial LLM providers when employees use tools like ChatGPT for work tasks, with no control over where that information goes. This document classifies every Phase 1 use case by risk using two dimensions, so the highest exposure use cases are identified and given the right level of care.

2 Categories

Every use case is placed into one of six categories, based on the kind of work it does rather than the topic it touches, so it can be scanned and compared consistently across the document.

- Coding
- QA and Testing
- Documentation
- Data Sorting and Analysis
- Rephrasing and Drafting
- Security and Risk Review

3 Risk Dimensions

3.1 Impact

Impact signifies what we gave the model. It reflects the sensitivity of the input itself, and whether that sensitivity is compounded by the fact that there is currently no generic, controlled platform in place, meaning inputs go to tools that retain and may train on them by default, with no control over how that data is used afterward.

Level	What it means	Why it gets this level
Low	The input has nothing in it beyond what could be shared freely with anyone outside the company.	Even though the tool may retain the input, there is nothing sensitive being retained, so there is nothing here that could harm the business, even in the worst case.
Medium	The input has real internal business information that is not public, but is also not uniquely valuable, such as internal reports or process details.	If this is retained or used for training with no control over its future use, it would cause real but recoverable trouble, such as embarrassment or friction, not lasting damage.
High	The input has proprietary or private information, such as source code, system design, trade secrets, or another company's confidential data.	Because there is no controlled platform, this input may be retained or trained on indefinitely. This is not a one time exposure, it is a permanent loss of control over that information, and it could seriously harm the business, hurt a client relationship, or create legal trouble that cannot easily be undone.

3.2 Fallout

Fallout tells us what happens if the model gives a wrong answer and nobody catches it. Unlike Impact, this has nothing to do with what was shared, it is only about what happens after the answer is given, so a use case with completely harmless input can still cause real damage if a confident but wrong answer gets trusted and acted on without anyone reviewing it first.

Level	What it means	Why it gets this level
Low	A wrong answer is either obvious right away, or it does not lead to anything serious even if nobody catches it.	There is no real path from a wrong answer to actual harm. The mistake shows itself, or it simply does not matter enough to act on.
Medium	A wrong answer causes wasted time or extra work, but someone eventually catches it before it goes further.	The cost is real, but it stays contained. It costs effort, not money, safety, or trust.
High	A wrong answer could be acted on before anyone checks it, and that action leads to a real financial, security, legal, or client problem.	This is the case where a wrong answer, if nobody catches it, causes real and sometimes lasting harm, such as a bad decision, a leaked issue, or a legal problem.

4 Overall Risk

Tier	Combination
Tier 1	High in Impact and High in Fallout
Tier 2	High in Impact and Medium in Fallout
Tier 3	Medium in Impact and Medium in Fallout
Tier 4	Medium in Impact and Low in Fallout
Tier 5	Low in Impact and Low in Fallout

5 Use Case Risk Matrix

Coding

Use case	Impact	Fallout	Tier	Justification
Writing code	High	Medium	Tier 2	Needs access to real company code. Bugs are usually caught in review before release.
Architecture and design brainstorming	High	Medium	Tier 2	Touches real system design and technical decisions specific to the company. Engineers evaluate and adapt suggestions before acting on them, so a flawed idea typically gets refined out during review, though it can still cost real rework once implementation is underway if a wrong assumption slips through.
Debugging code	High	Medium	Tier 2	Needs access to real code. Wrong fixes are usually caught during testing before release.
Code review	High	High	Tier 1	Sees real code. A missed issue here has no other check before the code is merged.
Writing and explaining SQL queries	Medium	Medium	Tier 3	Shows the shape of the database. A wrong query can quietly give a misleading result if nobody checks it.

QA and Testing

Use case	Impact	Fallout	Tier	Justification
Root cause analysis assistance	Medium	High	Tier 2	Uses internal system information. A wrong diagnosis can sound convincing and get acted on before anyone checks.
Troubleshooting steps	Medium	Medium	Tier 3	Uses internal system information. Wrong steps are usually noticed when the problem is still not fixed.
Test case and test plan generation	Medium	Medium	Tier 3	Sees feature details. Gaps in coverage may not show up until later testing.
Generating test cases	Medium	Medium	Tier 3	Sees feature details. Missing cases can let a real problem slip through if nobody reviews it.
Bug report writing	Medium	Low	Tier 4	Describes internal product behaviour. Mistakes are usually caught during triage.

Documentation

Use case	Impact	Fallout	Tier	Justification
API documentation	High	Medium	Tier 2	Describes real system structure. Misleading documentation is usually caught when something fails to connect properly.
Writing user manuals and product docs	Medium	Medium	Tier 3	Describes product details. A mistake that reaches a published manual could mislead an outside reader.
Drafting SOPs and process docs	Medium	Medium	Tier 3	Describes an internal process. An unreviewed mistake can become a permanent wrong step.
Internal wiki and knowledge base articles	Medium	Medium	Tier 3	Becomes a lasting internal reference. An uncaught mistake can spread to everyone who reads it.
Slides and presentation content	Medium	Low	Tier 4	May include internal business information. The presenter usually checks it before sharing.
Policy drafting	Medium	Medium	Tier 3	Internal governance content. An unclear clause can cause confusion until someone reviews it.
Job description drafting	Low	Low	Tier 5	Fairly general content, and HR checks it before posting. Low harm even if something is off.
Onboarding docs	Low	Low	Tier 5	General process information, checked by a manager or HR before use.
FAQ and help centre content	Low	Low	Tier 5	Meant to be public anyway, and checked before it goes live. Low harm even if slightly off.

Data Sorting and Analysis

Use case	Impact	Fallout	Tier	Justification
Spreadsheet and data analysis	Medium	High	Tier 2	Works with internal business data. A wrong number can feed directly into a decision before anyone rechecks it.
Budget and expense analysis	Medium	High	Tier 2	Works with internal financial data. An unnoticed mistake can shape a real budget decision.
Forecasting and financial modelling	Medium	High	Tier 2	Uses internal financial data. A flawed forecast can drive a decision before it is checked independently.
Resume screening and summarising	Medium	High	Tier 2	Involves a candidate's personal details. A wrong summary can affect a hiring decision if nobody double checks it.
Summarising reports and data	Medium	Medium	Tier 3	Summarises internal reports. A wrong figure can mislead a reader who does not check the original.
Market and competitor research	Low	Medium	Tier 4	Uses information that is already public. A confident but wrong conclusion could still steer a decision.
Ticket triage and categorisation	Medium	Low	Tier 4	Works with customer support content. Wrong categorising is usually noticed and fixed quickly.

Rephrasing and Drafting

Use case	Impact	Fallout	Tier	Justification
Summarising long documents and threads	Medium	Low	Tier 4	May include internal discussion. The original document is still there to check against.
Meeting notes and minutes	Medium	Low	Tier 4	May capture internal discussion. Attendees can catch anything wrong against what they remember.
Drafting customer and client replies	Medium	Medium	Tier 3	May mention a customer's account details. An unreviewed wrong reply can affect the relationship.
Performance review drafting	Medium	Medium	Tier 3	Involves an employee's performance information. An inaccurate draft could affect someone if it goes uncorrected.
Vendor communication drafting	Medium	Medium	Tier 3	May mention contract or vendor terms. An unreviewed mistake could affect the vendor relationship.
Invoice and financial report drafting	Medium	Medium	Tier 3	Involves internal financial figures. Mistakes are usually caught during normal finance review.
Proofreading and editing	Low	Low	Tier 5	No real judgment involved beyond fixing a draft. Mistakes are obvious right away.
Drafting emails	Low	Low	Tier 5	Ordinary content with no real judgment involved. The sender checks it before sending.
Translating text	Low	Low	Tier 5	A simple transformation task. A bilingual reviewer can catch mistakes before use.
Scheduling and admin coordination	Low	Low	Tier 5	Routine admin content with no real judgment involved. Mistakes are noticed and fixed quickly.

Security and Risk Review

Use case	Impact	Fallout	Tier	Justification
Code vulnerability review	High	High	Tier 1	Sees real source code. A missed vulnerability means it ships without anyone knowing.
Analysing logs and alerts	High	High	Tier 1	Sees internal system and security logs. A missed alert can let a real security problem go unnoticed.
Threat and CVE research	High	High	Tier 1	Tied directly to the company's security posture. A missed threat can leave a real weakness unaddressed.
Contract and policy review	High	High	Tier 1	Involves private contract terms. A misread clause can create legal trouble before anyone catches it.
Incident report drafting	High	Medium	Tier 2	Documents sensitive incident details. Security reviews it before action, but the detail itself is still very sensitive.

Totals: Tier 1: 5, Tier 2: 10, Tier 3: 13, Tier 4: 6, Tier 5: 7

6 What Tier 1 and Tier 2 Are Actually Exposing

The 15 Tier 1 and Tier 2 use cases — the two tiers where Impact hits High — are not all dangerous for the same reason. This section groups them by what they actually expose, so the pattern is visible rather than buried in 15 separate rows.

Exposure Type	Count	Use Cases
IP and proprietary	6	Writing code, Architecture and design brainstorming, Debugging code, Code review, API documentation, Code vulnerability review
Security posture	3	Analysing logs and alerts, Threat and CVE research, Incident report drafting
Financial data	3	Spreadsheet and data analysis, Budget and expense analysis, Forecasting and financial modelling
Legal and contractual	1	Contract and policy review
Personal data	1	Resume screening and summarising
Internal operational	1	Root cause analysis assistance

6 of the 15 Tier 1 and Tier 2 use cases expose IP or proprietary information. That is the single largest group in the two highest tiers, and it matches the priority already set for this project. IP exposure is treated as more serious than financial exposure, and the numbers here show why that priority matters in practice, not just in principle.

7 Use Cases by Tier

Tier 1 — 5 use cases

Use case	Category	Impact	Fallout
Code review	Coding	High	High
Code vulnerability review	Security and Risk Review	High	High
Analysing logs and alerts	Security and Risk Review	High	High
Threat and CVE research	Security and Risk Review	High	High
Contract and policy review	Security and Risk Review	High	High

Tier 2 – 10 use cases

Use case	Category	Impact	Fallout
Writing code	Coding	High	Medium
Architecture and design brainstorming	Coding	High	Medium
Debugging code	Coding	High	Medium
Root cause analysis assistance	QA and Testing	Medium	High
API documentation	Documentation	High	Medium
Spreadsheet and data analysis	Data Sorting and Analysis	Medium	High
Budget and expense analysis	Data Sorting and Analysis	Medium	High
Forecasting and financial modelling	Data Sorting and Analysis	Medium	High
Resume screening and summarising	Data Sorting and Analysis	Medium	High
Incident report drafting	Security and Risk Review	High	Medium

Tier 3 – 13 use cases

Use case	Category	Impact	Fallout
Writing and explaining SQL queries	Coding	Medium	Medium
Troubleshooting steps	QA and Testing	Medium	Medium
Test case and test plan generation	QA and Testing	Medium	Medium
Generating test cases	QA and Testing	Medium	Medium
Writing user manuals and product docs	Documentation	Medium	Medium
Drafting SOPs and process docs	Documentation	Medium	Medium
Internal wiki and knowledge base articles	Documentation	Medium	Medium
Policy drafting	Documentation	Medium	Medium
Summarising reports and data	Data Sorting and Analysis	Medium	Medium
Drafting customer and client replies	Rephrasing and Drafting	Medium	Medium
Performance review drafting	Rephrasing and Drafting	Medium	Medium
Vendor communication drafting	Rephrasing and Drafting	Medium	Medium
Invoice and financial report drafting	Rephrasing and Drafting	Medium	Medium

Tier 4 — 6 use cases

Use case	Category	Impact	Fallout
Bug report writing	QA and Testing	Medium	Low
Slides and presentation content	Documentation	Medium	Low
Market and competitor research	Data Sorting and Analysis	Low	Medium
Ticket triage and categorisation	Data Sorting and Analysis	Medium	Low
Summarising long documents and threads	Rephrasing and Drafting	Medium	Low
Meeting notes and minutes	Rephrasing and Drafting	Medium	Low

Tier 5 — 7 use cases

Use case	Category	Impact	Fallout
Job description drafting	Documentation	Low	Low
Onboarding docs	Documentation	Low	Low
FAQ and help centre content	Documentation	Low	Low
Proofreading and editing	Rephrasing and Drafting	Low	Low
Drafting emails	Rephrasing and Drafting	Low	Low
Translating text	Rephrasing and Drafting	Low	Low
Scheduling and admin coordination	Rephrasing and Drafting	Low	Low

8 Use Case Buckets

Coding

Use Case	Adoption	Tier	Model Category
Writing code	48%	Tier 2	Coding
Architecture and design brainstorming	42%	Tier 2	Coding
Debugging code	38%	Tier 2	Coding
Code review	35%	Tier 1	Coding
Writing and explaining SQL queries	15%	Tier 3	Coding

QA and Testing

Use Case	Adoption	Tier	Model Category
Root cause analysis assistance	33%	Tier 2	Deep
Troubleshooting steps	31%	Tier 3	Deep
Test case and test plan generation	29%	Tier 3	General
Generating test cases	21%	Tier 3	Coding
Bug report writing	15%	Tier 4	Quick

Documentation

Use Case	Adoption	Tier	Model Category
API documentation	19%	Tier 2	Coding
Writing user manuals and product docs	44%	Tier 3	General
Drafting SOPs and process docs	42%	Tier 3	General
Internal wiki and knowledge base articles	35%	Tier 3	General
Slides and presentation content	42%	Tier 4	General
Policy drafting	8%	Tier 3	General
Job description drafting	21%	Tier 5	Quick
Onboarding docs	6%	Tier 5	General
FAQ and help centre content	4%	Tier 5	Quick

Data Sorting and Analysis

Use Case	Adoption	Tier	Model Category
Spreadsheet and data analysis	38%	Tier 2	Deep
Budget and expense analysis	17%	Tier 2	Deep
Forecasting and financial modelling	12%	Tier 2	Deep
Resume screening and summarising	6%	Tier 2	General
Summarising reports and data	65%	Tier 3	General
Market and competitor research	40%	Tier 4	General
Ticket triage and categorisation	12%	Tier 4	Quick

Rephrasing and Drafting

Use Case	Adoption	Tier	Model Category
Summarising long documents and threads	65%	Tier 4	General
Meeting notes and minutes	23%	Tier 4	Quick
Drafting customer and client replies	19%	Tier 3	Quick
Performance review drafting	15%	Tier 3	General
Vendor communication drafting	10%	Tier 3	Quick
Invoice and financial report drafting	6%	Tier 3	General
Proofreading and editing	56%	Tier 5	Quick
Drafting emails	50%	Tier 5	Quick
Translating text	15%	Tier 5	Quick
Scheduling and admin coordination	2%	Tier 5	Quick

Security and Risk Review

Use Case	Adoption	Tier	Model Category
Code vulnerability review	29%	Tier 1	Coding
Analysing logs and alerts	17%	Tier 1	Deep
Threat and CVE research	2%	Tier 1	Deep
Contract and policy review	2%	Tier 1	Deep
Incident report drafting	15%	Tier 2	General

9 Model Recommendations

Coding: to be decided, no candidates researched yet.

Quick Assistant

Model	Role	VRAM (Q4)	License	Note
Mistral 7B	Incumbent	4.4 GB	Apache 2.0	Locally tested, number 2 overall. Best summary fidelity and length discipline.
Qwen3.5-9B	Candidate	5.5 GB	Apache 2.0	Strongest model under 10B parameters on general reasoning benchmarks. Not yet tested.
Gemma 4 12B	Candidate	6.6 GB	Apache 2.0	Beats Gemma 3 27B on MMLU Pro at less than half the size.

Model	Role	VRAM (Q4)	License	Note
Minstral 3 8B	Candidate	4.8 GB	Apache 2.0	Mistral's own successor to the tested 7B. Lowest risk fallback.

General Assistant

Model	Role	VRAM (Q4)	License	Note
GPT-OSS 20B	Incumbent	~13 GB	Apache style (OpenAI)	Never empirically tested. Carrying majority of daily work on spec sheet trust alone.
Qwen3.6-27B	Candidate	16.8 GB	Apache 2.0	Ties Sonnet 4.6 on AA Agentic Index. Untested locally.
NVIDIA Nemotron 3 Nano	Candidate	19.0 GB	NVIDIA open license	Ahead of GPT-OSS 20B and Qwen3-30B-A3B on accuracy, more than 2x throughput.
GLM-4.7-Flash	Candidate	18.0 GB	Custom (verify terms)	Runs on consumer GPU. Similar efficiency profile to Nemotron 3 Nano.

Deep Assistant

Model	Role	VRAM (Q4)	License	Note
DeepSeek R1 Distill Qwen 32B	Incumbent	19.5 GB	MIT	Best local reasoning after baseline. Needs 3 configs and a pinned seed to terminate. Rumination bug confirmed at 8B, untested at 32B.
Qwen3-32B (thinking mode)	Candidate	19.7 GB	Apache 2.0	Trained natively through reinforcement learning. Beats DeepSeek R1 Distill Llama 70B at half the size.
NVIDIA Nemotron 3 Nano Reasoning	Candidate	18.6 GB	NVIDIA open license	Dedicated reasoning variant, NVFP4 native.
QwQ-32B	Candidate	19.5 GB	Apache 2.0	Most proven reasoning model on the list, with more production hours behind it than other candidates. Matches DeepSeek R1 671B on math and coding at a fraction of the size.

Quantization

- Well calibrated Q4 (AWQ, GPTQ, Q4_K_M) against an FP16 baseline typically costs about 1 to 3% on general tasks, drafting, summarizing, and everyday writing.
- Complex multi step reasoning degrades more at the same bit width, roughly 3 to 5%, about double the general task hit, per one benchmark source.
- FP16 uses 2 bytes per param, matching every reported figure exactly. Q4 uses about 0.60 bytes per param, calibrated against four confirmed Q4_K_M sizes.

- MoE rows are sized on total params, since every expert must be VRAM resident even though only the active fraction computes per token.

Assistant	Model	Params	FP16	FP8	Q4
Quick	Mistral 7B	7.3B dense	14.6 GB	7.7 GB	4.4 GB (confirmed)
Quick	Qwen3.5-9B	9B dense	18.0 GB	9.5 GB	5.5 GB (confirmed)
Quick	Gemma 4 12B	11.95B dense	23.9 GB	12.5 GB	6.6 GB (confirmed)
Quick	Ministral 3 8B	8B dense	16.0 GB	8.4 GB	4.8 GB (est.)
General	GPT-OSS 20B	20.9B / 3.6B active	—	—	~13 GB (ships native)
General	Qwen3.6-27B	27.8B dense	55.6 GB	29.2 GB	16.8 GB (confirmed)
General	Nemotron 3 Nano	31.6B / 3.2B active	63.2 GB	33.2 GB	19.0 GB (est.)
General	GLM-4.7-Flash	30B / 3B active	60.0 GB	31.5 GB	18.0 GB (est.)
Deep	DeepSeek R1 Distill 32B	32.5B dense	65.0 GB	34.1 GB	19.5 GB (est.)
Deep	Qwen3-32B	32.8B dense	65.6 GB	34.4 GB	19.7 GB (confirmed)
Deep	Nemotron 3 Nano Reasoning	31B / 3B active	62.0 GB	32.6 GB	18.6 GB (est.)
Deep	QwQ-32B	32.5B dense	65.0 GB	34.1 GB	19.5 GB (est.)



Use Case Scoring Rationale

Companion to Use Case Risk Recategorization

Ryef Taimur, Ibrahim Murtaza

Purpose

Use Case Risk Recategorization scores every Phase 1 use case on two dimensions, Impact and Fallout, and places it into one of five tiers. This document is the supporting rationale: for each use case, it explains separately why it received its Impact score and why it received its Fallout score, rather than collapsing both into a single justification line. Category order and use case names match the parent document exactly, so the two can be read side by side.

Impact asks what we gave the model, reflecting both the sensitivity of the input and the fact that there is currently no controlled platform in place, so what is shared may be retained or trained on with no control over its future use. Fallout asks what happens if the model's answer is wrong and nobody notices. The two are independent: a use case can expose sensitive input and still carry low downstream risk if a human always checks the output before acting on it, or the reverse, if the input is harmless but an unchecked wrong answer could still cause damage.

Coding

Writing code

Tier 2

High

Impact

Requires sharing real company source code with the model — the actual proprietary implementation, not a description of it.

Medium

Fallout

Bugs in generated code are normally caught in review or testing before release, so a wrong answer costs rework rather than reaching production.

Architecture and design brainstorming

Tier 2

High

Impact

Touches real system design and technical decisions specific to the company. Architecture choices count as proprietary even when no code is shared.

Medium

Fallout

Engineers evaluate and adapt suggestions before acting on them, so a flawed idea typically gets refined out during review, though it can still cost real rework once implementation is underway if a wrong assumption slips through.

Debugging code

Tier 2

High

Impact

Requires sharing real code plus internal error context, such as logs or stack traces, to diagnose the problem.

Medium

Fallout

A wrong fix is normally caught during testing before release, so the failure mode is wasted effort, not a shipped defect.

Code review

Tier 1

High

Impact

Shares real, complete code with the model for review — the full proprietary artifact, not a fragment.

High

Fallout

Review is usually the last check before merge. If the model misses a real issue, there is no downstream safety net catching it before the code ships.

Writing and explaining SQL queries**Tier 3**

- Medium Impact** Reveals the shape of the database, such as table and column names, without exposing the proprietary business logic built on top of it.
- Medium Fallout** A subtly wrong query can return a plausible but misleading result. Someone usually notices when the output does not match expectations, but not always immediately.

QA and Testing**Root cause analysis assistance****Tier 2**

- Medium Impact** Uses internal system information, such as logs, symptoms, and configuration details, but not source code or trade secrets directly.
- High Fallout** A wrong diagnosis can sound just as confident as a correct one and get acted on. Teams may chase the wrong fix before anyone double-checks the reasoning.

Troubleshooting steps**Tier 3**

- Medium Impact** Uses internal system information at a process level, not deep proprietary detail.
- Medium Fallout** Wrong steps are usually noticed quickly, because the underlying problem is still visibly unresolved after following them.

Test case and test plan generation**Tier 3**

- Medium Impact** Sees feature details, not full source code or system design.
- Medium Fallout** A gap in test coverage does not fail loudly. It can sit unnoticed until a real defect the missing case would have caught reaches a later stage.

Generating test cases**Tier 3**

- Medium Impact** Sees feature details, not full source code or system design.
- Medium Fallout** Missing edge cases silently reduce coverage. The risk only shows up if a bug slips through that a more complete set would have caught.

Bug report writing**Tier 4**

- Medium Impact** Describes internal product behaviour, not proprietary implementation detail.
- Low Fallout** Bug reports get triaged by a human before any action is taken, so an inaccurate report is caught early and cheaply.

Documentation

API documentation

Tier 2

- High Impact** Describes real system structure, endpoints, and internal architecture — effectively documenting how the proprietary system is built.
- Medium Fallout** Misleading documentation tends to surface fast, because a developer following it hits a failure when something does not connect as described.

Writing user manuals and product docs

Tier 3

- Medium Impact** Describes product details, but not the underlying proprietary design.
- Medium Fallout** A mistake that reaches a published manual could mislead an external reader before anyone catches it.

Drafting SOPs and process docs

Tier 3

- Medium Impact** Describes an internal process, not proprietary technical detail.
- Medium Fallout** An unreviewed mistake can quietly become the de facto correct process if nobody catches it before publication.

Internal wiki and knowledge base articles

Tier 3

- Medium Impact** Contains internal reference content, not uniquely sensitive material.
- Medium Fallout** Becomes a lasting reference. An uncaught mistake keeps getting read and repeated by whoever consults it.

Slides and presentation content

Tier 4

- Medium Impact** May include internal business information in summarised form.
- Low Fallout** The presenter reviews the deck before sharing it, which catches most errors before they reach an audience.

Policy drafting

Tier 3

- Medium Impact** Contains internal governance content, not technical secrets.
- Medium Fallout** An unclear or wrong clause can cause confusion or inconsistent application until someone reviews and corrects it.

Job description drafting**Tier 5****Low Impact** Fairly generic content, not tied to internal systems or confidential data.**Low Fallout** HR reviews postings before they go live, and the cost of an off draft is minor.**Onboarding docs****Tier 5****Low Impact** General process information with nothing proprietary in it.**Low Fallout** Checked by a manager or HR before use; low stakes even if imperfect.**FAQ and help centre content****Tier 5****Low Impact** Meant to be public anyway.**Low Fallout** Checked before it goes live; the harm from a small error is minimal.**Data Sorting and Analysis****Spreadsheet and data analysis****Tier 2****Medium Impact** Works with internal business data that is not uniquely valuable on its own.**High Fallout** A wrong number can feed directly into a decision before anyone rechecks the underlying analysis.**Budget and expense analysis****Tier 2****Medium Impact** Works with internal financial data, not trade secrets.**High Fallout** An unnoticed mistake can directly shape a real budget decision.**Forecasting and financial modelling****Tier 2****Medium Impact** Uses internal financial data at a business-planning level.**High Fallout** A flawed forecast can drive a decision before it is checked independently.

Resume screening and summarising**Tier 2**

- Medium Impact** Involves a candidate's personal details — real, but not uniquely damaging on their own.
- High Fallout** A wrong summary can affect a hiring decision if nobody double-checks it against the original resume.

Summarising reports and data**Tier 3**

- Medium Impact** Summarises internal reports that are not public.
- Medium Fallout** A wrong figure can mislead a reader who trusts the summary rather than checking the source.

Market and competitor research**Tier 4**

- Low Impact** Uses information that is already public.
- Medium Fallout** A confidently wrong conclusion can still steer a decision, even though the underlying inputs were public.

Ticket triage and categorisation**Tier 4**

- Medium Impact** Works with customer support content that is internal but not uniquely sensitive.
- Low Fallout** Miscategorised tickets are usually noticed and corrected quickly during normal handling.

Rephrasing and Drafting**Summarising long documents and threads****Tier 4**

- Medium Impact** May include internal discussion that is not public.
- Low Fallout** The original document is still available to check against, so errors are easy to catch.

Meeting notes and minutes**Tier 4**

- Medium Impact** May capture internal discussion.
- Low Fallout** Attendees can catch anything wrong against their own memory of the meeting.

Drafting customer and client replies

Tier 3

Medium Impact May mention a customer's account details.**Medium Fallout** An unreviewed wrong reply can affect the customer relationship before anyone catches it.**Performance review drafting**

Tier 3

Medium Impact Involves an employee's personal performance information.**Medium Fallout** An inaccurate draft could unfairly affect someone if it goes uncorrected into a real review.**Vendor communication drafting**

Tier 3

Medium Impact May mention contract or vendor terms.**Medium Fallout** An unreviewed mistake could affect the vendor relationship or misstate a commitment.**Invoice and financial report drafting**

Tier 3

Medium Impact Involves internal financial figures.**Medium Fallout** Mistakes are usually caught during normal finance review before anything is sent externally.**Proofreading and editing**

Tier 5

Low Impact No sensitive content beyond what is already in the draft being fixed.**Low Fallout** Mistakes are obvious right away, since the task itself is mechanical.**Drafting emails**

Tier 5

Low Impact Ordinary content with nothing proprietary in it.**Low Fallout** The sender checks it before sending.

Translating text

Tier 5

Low Impact A simple transformation task with no added proprietary information.

Low Fallout A bilingual reviewer can catch mistakes before use.

Scheduling and admin coordination

Tier 5

Low Impact Routine admin content with nothing sensitive in it.

Low Fallout Mistakes are noticed and fixed quickly, since scheduling conflicts surface fast.

Security and Risk Review**Code vulnerability review**

Tier 1

High Impact Sees real source code in full.

High Fallout A missed vulnerability ships silently — there is no other check standing between this review and production.

Analysing logs and alerts

Tier 1

High Impact Sees internal system and security logs directly.

High Fallout A missed alert can let a real security problem go unnoticed until it escalates.

Threat and CVE research

Tier 1

High Impact Tied directly to the company's actual security posture.

High Fallout A missed or misjudged threat can leave a real weakness unaddressed.

Contract and policy review

Tier 1

High Impact Involves private contract terms and confidential clauses.

High Fallout A misread clause can create legal exposure before anyone catches the error.

Incident report drafting

Tier 2

High	Impact	Documents sensitive, real incident detail.
Medium	Fallout	Security reviews the report before any action is taken on it, so a mistake is caught before it drives a decision, though the underlying detail is still highly sensitive regardless of that review step.